

AI Command Center

Lets seller talk to the platform: ask questions, create drafts, update safe settings, and perform approved actions.

Why this exists

This makes the platform feel powerful without forcing sellers through many menus.

Primary user	Owner and authorized staff.
Business value	Differentiator. Saves time and makes advanced features easier to use.
Build priority	P1 - after core data/actions exist.

1. Page Goal

- Provide global AI command bar and dedicated AI page.
- Answer business questions from orders, products, inventory, ads, and customers.
- Draft actions and require confirmation for risky work.
- Show exact preview before executing changes.

2. How It Should Look

Global command bar: 'Ask AI...' with suggestions by current page.
AI page: chat history, suggested tasks, pending approvals, recent actions.
Result cards: answer, chart/table, draft action, confirm button.
Audit panel: what AI did, who approved, when.

3. Main Sections

Ask data	Today sales, pending orders, low stock, best product.
Draft work	Create coupon, write campaign, prepare product description.
Execute safe action	Change draft text, tag customers, create report.
Approval flow	Confirm before courier, payment, bulk message, delete.
Audit	History of AI suggestions and actions.

4. User Actions

Ask question	Returns answer from shop data.	If data unavailable, say what is missing.
Create draft	Builds draft campaign/coupon/product text.	Seller reviews before publish.

Run action	Executes approved safe command.	Risky command shows preview and confirm.
Undo	Reverts supported AI action.	If not undoable, explain before running.
View source	Shows data used for answer.	Never expose another shop's data.

5. Data Fields

- Command: user, text, page context, intent, risk level, result.
- Approval: approver, timestamp, before/after preview, action status.
- Audit: source data, tool/action used, error, rollback status.
- Permissions: which roles can ask, draft, approve, execute.

6. AI / Automation Behavior

- AI must classify command risk before acting.
- AI can answer and draft freely; execution depends on permission/risk.
- AI must cite shop data source for important business answers.
- AI refuses actions outside connected shop or permissions.

7. Empty, Error, and Safety States

- Empty state must show one clear next button, not a blank table.
- Every destructive action needs confirm: delete, cancel order, disconnect account, refund mark.
- If third-party API fails, keep local draft and show retry.
- AI must never invent price, stock, delivery charge, delivery date, or policy.

8. Acceptance Checklist

- Risky actions need preview and explicit approve.
- AI cannot bypass staff role permissions.
- Every executed command writes audit log.
- Bad command or unclear instruction asks one clarifying question.

9. Do Not Build Yet

- No fully autonomous business manager in V1.
- No silent background bulk changes.
- No cross-shop benchmarking until privacy/legal is solved.